

Binary Search Trees

COMP2611: Data Structures
2019/2020

Outline

- ▶ Notion of ordered retrieval
- ▶ Binary Search Trees
 - ▶ Definitions
 - ▶ CRUD operations - insert, search (for specific key), delete
 - ▶ Problems and a brief overview of the solutions to fixing them (Red-Black Trees and AVL Trees)

Ordered Retrieval

- ▶ Suppose that our data type of our keys have an ordering defined on them
- ▶ Want to retrieve keys and/or values sorted by keys
- ▶ Want to retrieve smallest or largest key (and maybe associated value if any)
- ▶ Can't do this using Hashtables
 - ▶ Why?

Naive Implementation

- ▶ Store in array!
- ▶ Use binary search to find data and to find place to insert order to retain ordering
- ▶ Deletion: search for key, remove data, move rest of array up
- ▶ Problems?
 - ▶ Search: $O(\log n)$:-)
 - ▶ Insert: $O(n)$:-)
 - ▶ Delete: $O(n)$:-)

Ordered Retrieval

- ▶ Important operation!
- ▶ Backbone of most RDMS!
- ▶ Useful when we temporal data
 - ▶ Can find out information such as “What even occurred on or before 25th September 2019?”
- ▶ Need efficient data structure and algorithm
- ▶ Solution: Binary Search Trees! (BSTs)


Binary Search Trees

- ▶ Special type of Binary Tree
- ▶ Each node has left field, right field, parent field and data field as before, but:
 - ▶ data field has two subfields: key, and (optionally) value
 - ▶ invariant on left field and right field:
 - ▶ All keys in left subtree are less than key in data
 - ▶ All keys in right subtree are greater than key in data
 - ▶ Corollary: in-order traversal of BST can be used to print keys in ascending order!
 - ▶ Can prove this formally!

Binary Search Trees


- ▶ We need to think recursively!
- ▶ Useful tip: typically need to think about 5 cases:
 - ▶ When a tree is empty
 - ▶ When a node has no children
 - ▶ When a node has only a left child
 - ▶ When a node has only a right child
 - ▶ When a node has both a left and right child
- ▶ Often the action we take is common throughout many cases

Searching a BST



- ▶ Find **11**
- ▶ Each comparison tells us whether to go left or right

Searching a BST




- ▶ What if item isn't in tree?
- ▶ Find 14

Searching a BST

```
function search(root, key_to_find):  
    if root is NIL:  
        return NIL  
    if root.data.key == key_to_find:  
        return root  
    if root.data.key < key_to_find:  
        return search(root.left, key_to_find)  
    return search(root.right, key_to_find)
```

Inserting in a BST



- ▶ To insert, perform a search and add as new leaf
- ▶ Insert 17

Inserting into a BST

```
function insert(root, key, value=NIL):  
    if root.data.key == key:  
        return  
    if root.data.key < key:  
        if root.left is NIL:  
            data = create_data(key, value)  
            root.left = new_node(data, root)  
        else:  
            insert(root.left, key, value)  
    if root.data.key > key:  
        if root.right is NIL:  
            data = create_data(key, value)  
            root.right = new_node(data, root)  
        else:  
            insert(root.right, key, value)
```

Smallest Node

- ▶ Smallest node is the node with the smallest key in the BST
- ▶ How to find? Remember the cases we need to consider
 - ▶ Case #1: Tree is empty
 - ▶ Case #2: Root has no children (root is only node)
 - ▶ Case #3: Root has only a left child
 - ▶ Case #4: Root has only a right child
 - ▶ Case #5: Root has both left and right child

Smallest Node

- ▶ Smallest node is the node with the smallest key in the BST
- ▶ How to find? Remember the cases we need to consider
 - ▶ Case #1: Tree has no smallest node
 - ▶ Case #2: Root is smallest node
 - ▶ Case #3: Smallest node is in left child
 - ▶ Case #4: Root is smallest node
 - ▶ Case #5: Smallest node in left child

Smallest Node

- ▶ Lefts combine cases!
- ▶ Case# 1 stands alone
- ▶ Case#2 and Case#4 are the same (what is common between them?)
- ▶ Case #3 and Case#5 are the same (what is common between them?)

Smallest Node

- ▶ Lefts combine cases!
- ▶ Case#1 stands alone
- ▶ Case#2 and Case#4 are the same (no left child)
- ▶ Case #3 and Case#5 are the same (has a left child)
- ▶ Can write algorithm from this!

Smallest Node in BST

```
function smallest(root):  
    if root is None:  
        return NIL  
    if root.left is NIL:  
        return root  
    return smallest(root.left)
```


Should try solving for largest node!

Bounded Range Retrieval

- ▶ We have an upper bound x and a lower bound y
- ▶ Want to retrieve all nodes in a BST with keys greater than or equal to x and less than or equal to y
- ▶ Want to retrieve in order of key
 - ▶ First, we find root of subtree that is within range (how would we do this?)
 - ▶ Then use a modification of in-order traversal!

Bounded Range Retrieval

- ▶ Case #1: If the root is empty, nothing to retrieve
- ▶ Case #2:
 - ▶ 2a: if the root has no children and is in the range, add it to retrieved items
 - ▶ 2b: if the root has no children and is in not range, terminate
- ▶ Case #3:
 - ▶ 3a: if the root has a left children and is in range, add it to retrieved items and go down left child
 - ▶ 3b: terminate

Bounded Range Retrieval

- ▶ Case #4:
 - ▶ 4a: if the root has a right children and is in range, add it to retrieved items and go down right child
 - ▶ 4b: terminate Case #2:
- ▶ Case #5:
 - ▶ 5a: if the root has both children and is in range, add it to retrieved items and go down both children
 - ▶ 5b: terminate

Range Retrieval

- ▶ Core idea:
 - ▶ Once subtree is found
 - ▶ If root is in range, add to items
 - ▶ Go down children and apply recursively
- ▶ Find subtree using modification of search

Range Retrieval

```
function range_ret_helper(root, hi, lo):  
    if root is None or root.data.key < lo  
        or root.data.key > hi:  
        return []  
  
    res = [root]  
    res_from_left = range_ret_helper(root.left, hi, lo)  
    res_from_right = range_ret_helper(root.right, hi, lo)  
    res = concat(res, res_from_right)  
    res = concat(res_from_left, res)  
    return res
```


Range Retrieval

```
function range_ret(root, hi, lo):  
    if root is None:  
        return []  
    if root.data.key >= lo and root.data.key <= hi:  
        return range_ret_helper(root, hi, lo)  
    if root.data.key < lo:  
        return range_ret_helper(root.right, hi, lo)  
    return range_ret_helper(root.left, hi, lo)
```

Successors and Predecessors

- ▶ In-order successor of a node x - the node that comes immediately after x in an in-order retrieval
- ▶ In-order predecessor of a node x - is the node that comes immediately after x in an in-order retrieval

Successor

- ▶ Core idea:
 - ▶ If we have a right child, then the successor is the smallest node in the right child
 - ▶ If we have no right child, then the successor is one of the nodes ancestors
 - ▶ Parent might be in successor if node is left of parent
 - ▶ If node not left of parent, then we need to find closest ancestor who is left of their parent. The parent of that ancestor is the in-order successor

Smallest Node in BST

```
function successor(root):  
    if root is None:  
        return NIL  
    if root.right is not NIL:  
        return smallest(root.right)  
    curr = root  
    p = root.parent  
    while p != NIL and p.right == curr:  
        curr = p  
        p = curr.parent  
    return p
```

Predecessor

- ▶ Predecessor is similar, but with some conditions changed
- ▶ Try writing on your own
- ▶ Will cover in labs

LTE queries

- ▶ We have some key \mathbf{k} , want to find node with key \mathbf{x} such that \mathbf{x} is the largest key in the tree such that $\mathbf{x} \leq \mathbf{k}$

LTE Queries

- ▶ Break down by cases!
 - ▶ In all cases, if root exists, then root is answer if root's key is the key we are concerned with
 - ▶ Empty tree: no answer
 - ▶ No children: root is either answer or not
 - ▶ Left child: if root's key is greater than k , then answer must be in left child; else not in this subtree
 - ▶ Right child: if root's key is less than k , then answer might be in right child; else answer is root

LTE Query


```
function lte_search(root, key):  
    if root is None:  
        return NIL  
    if root.data.key == key:  
        return root  
    if root.data.key < key:  
        res = lte_seaerch(root.right, key)  
        return (res is NIL) ? root:res  
    return lte_root(root.left, key)
```

GTE Query

- ▶ GTE query is similar
- ▶ Try writing on own
- ▶ Part of the assignment!


Removing from a BST

- ▶ Can be tricky
- ▶ Three cases to consider
 - ▶ Removing a leaf: easy, just do it
 - ▶ Removing internal node w/ **1** child (e.g., **15**)
 - ▶ Removing internal node w/ **2** children (e.g., **7**)



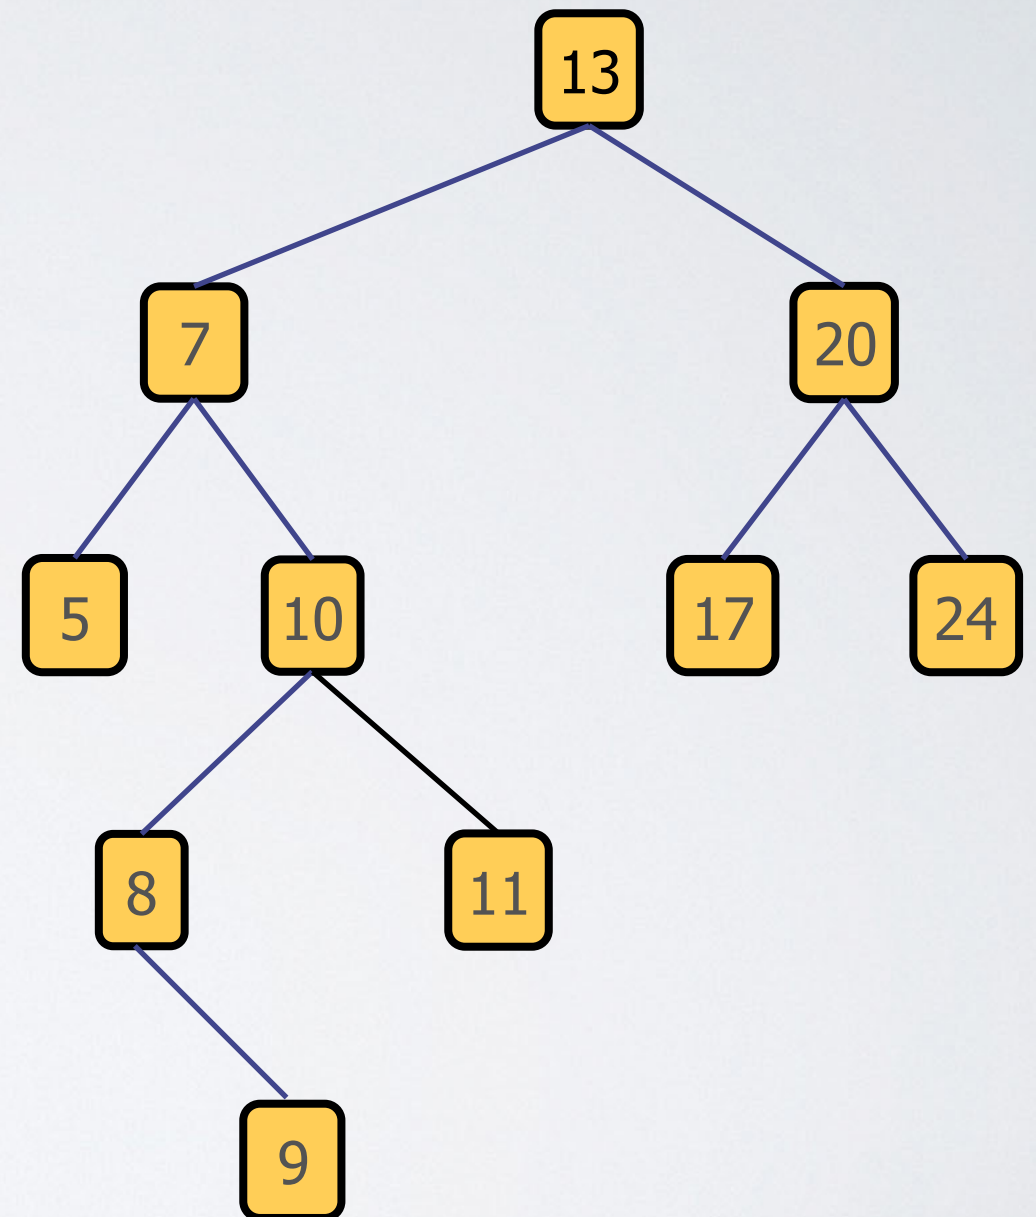
Removing from a BST - Case #2

- ▶ Removing internal node w/ 1 child
- ▶ Strategy
 - ▶ “Splice out” node by connecting its parent to its child
- ▶ Example: remove 15
 - ▶ set parent's left pointer to 17
 - ▶ remove 15's pointer
 - ▶ no more references to 15 so erased (garbage collected)
 - ▶ BST order is maintained



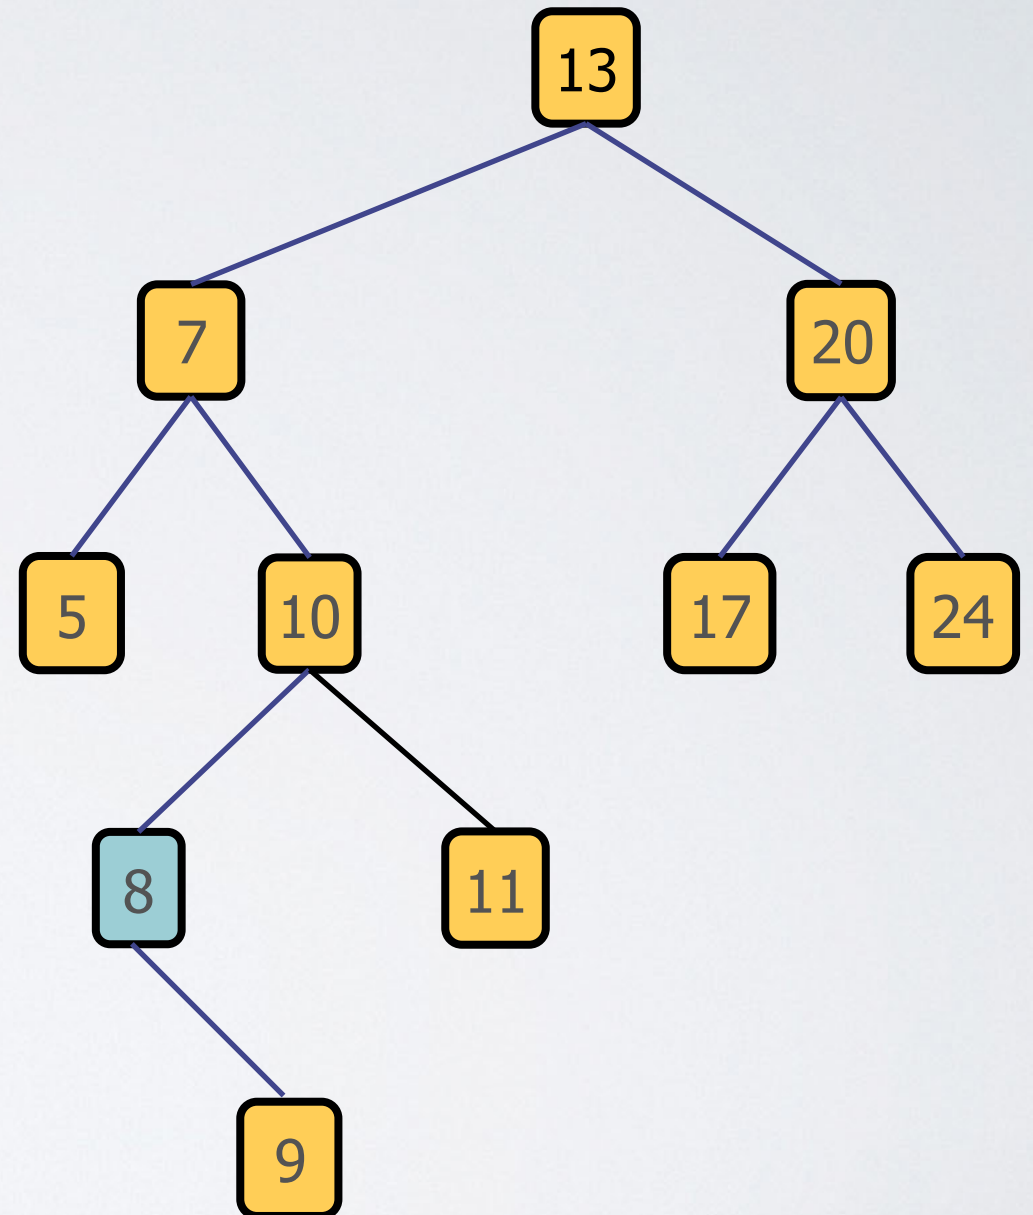
Removing from a BST - Case #3

- ▶ Removing internal node w/ **2** children
- ▶ Replace node w/ successor
 - ▶ successor: next largest node
- ▶ Delete successor
 - ▶ Successor a.k.a. the in-order successor
- ▶ Example: remove 7
 - ▶ What is successor of 7?



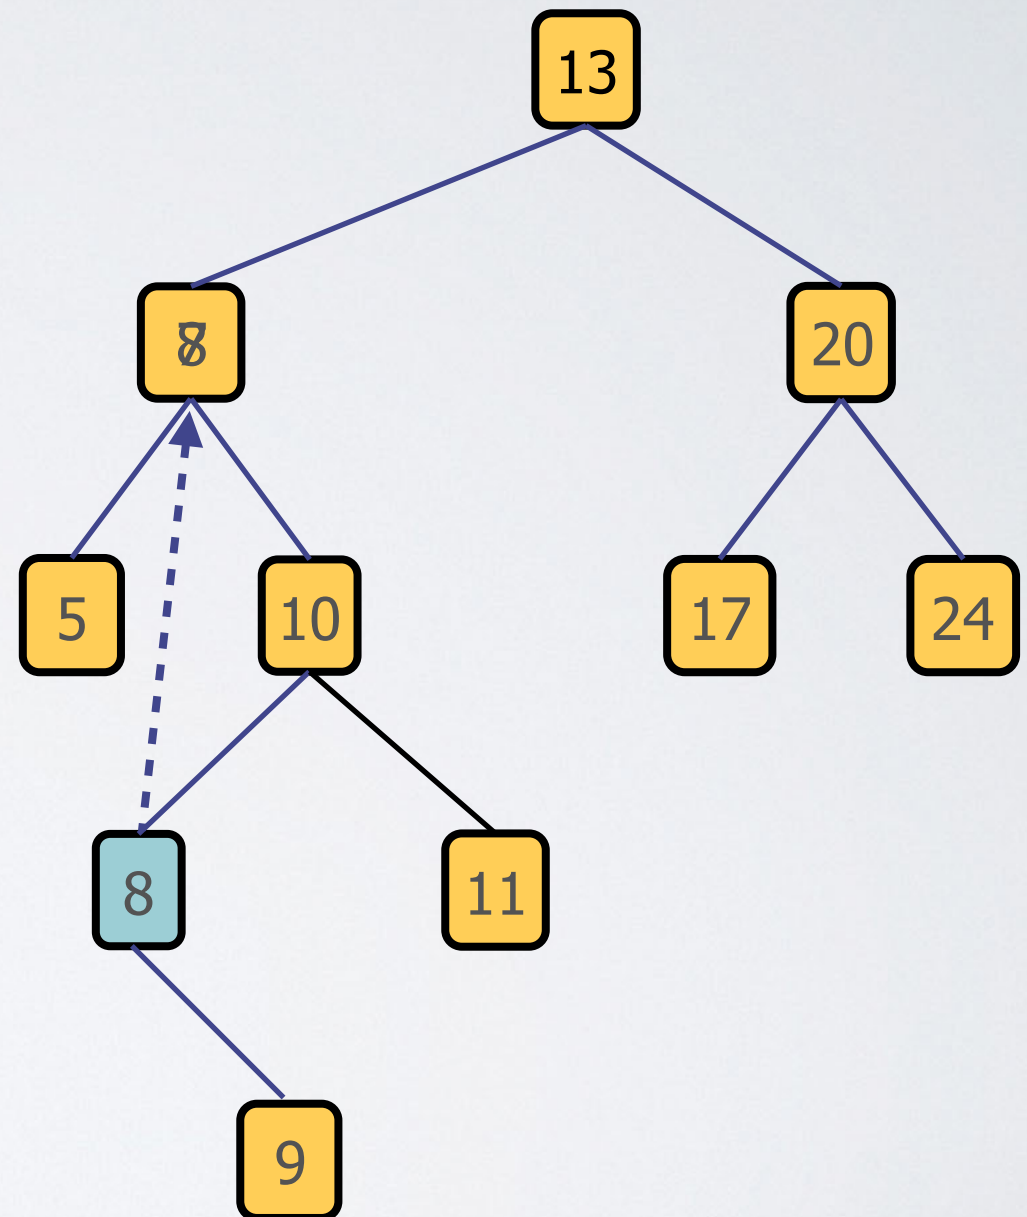
Removing from a BST - Case #3

- ▶ Since node has **2** children...
- ▶ ...it has a right subtree
- ▶ Successor is leftmost node in right subtree
- ▶ 7's successor is 8



Removing from a BST - Case #3

- ▶ Now, replace node with successor
- ▶ Observation
 - ▶ Successor can't have left sub-tree
 - ▶ ...otherwise its left child would be successor
 - ▶ so successor only has right child
- ▶ Remove successor using Case #1 or #2
 - ▶ Here, use case #2 (internal w/ 1 child)
- ▶ Successor removed and BST order restored




Deletion

- ▶ Try writing pseudocode as an exercise
- ▶ Will cover solution in lab


Binary Search Tree Analysis

- ▶ How fast are BST operations?
 - ▶ Given a tree, what is the worst-case node to find/remove?
- ▶ What is the best-case tree?
 - ▶ a balanced tree
- ▶ What is the worst-case tree?
 - ▶ a completely unbalanced tree



Degenerate Cases

- ▶ What if we insert sorted data (either ascending or descending)into BST?
- ▶ Tree looks like linked list
- ▶ Performance expectations breakdown!




How to solve?

- ▶ We modify our method of inserting data
- ▶ Use rotations to keep tree balanced or mostly balanced
- ▶ Can help guarantee $O(\log n)$ performance
- ▶ Techniques: AVL trees, Red-black trees

Binary Search Trees — Rotations

- ▶ We can re-balance unbalanced trees w/ tree rotations



- ▶ In-order traversal of all 3 trees is



- ▶ so BST order is preserved